

Project: Investigate a Movie Dataset (TMDb Movies Dataset)

Introduction

This is a project towards the Udcity Data Analytics nanodegree certification

This data set contains information about 10,000 movies collected from The Movie Database (TMDb), including user ratings and revenue.

- Certain columns, like 'cast' and 'genres', contain multiple values separated by pipe (|) characters.
- There are some odd characters in the 'cast' column. Don't worry about cleaning them. You can leave them as is.
- The final two columns ending with "_adj" show the budget and revenue of the associated movie in terms of 2010 dollars, accounting for inflation over time.

some questions of interest that arise from the movie features are:

- Which genres are most popular?
- How has movie runtime changed over the years?
- Which genres generate more revenue?
- How has overall movie revenue changed over the years?
- which is the most popular and least popular movie?
- What are the top ten most popular movies?
- What genres do most producers produce?
- How has the cost of movie production changed over the years?

However, we are going to answer the questions.

- How has The cost of movie production changed over the years?
- How has revenue from movie sales et'al changed over the years?
- How has profitability changed over the years?
- with respect to the 2010 inflation rate, comment on effect of onflation on overall cost overhead

we answer these questions using the actual revenue, values and the values adjusted for the 2010 value of the Dollar

```
#import the needed packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pprint as pp
import os
import sys
from code1 import *
```

```
%matplotlib inline
```

```
file = 'tmdb-movies.csv'
```

```
movie_df = pd.read_csv(file)
```

```
movie_df.head(3)
```

	id	imdb_id	popularity	budget	revenue	
original_title \						
0	135397	tt0369610	32.985763	150000000	1513528810	Jurassic World
1	76341	tt1392190	28.419936	150000000	378436354	Mad Max: Fury Road
2	262500	tt2908446	13.112507	110000000	295238201	Insurgent

	cast \
0	Chris Pratt Bryce Dallas Howard Irrfan Khan Vi...
1	Tom Hardy Charlize Theron Hugh Keays-Byrne Nic...
2	Shailene Woodley Theo James Kate Winslet Ansel...

	homepage	director \
0	http://www.jurassicworld.com/	Colin Trevorrow
1	http://www.madmaxmovie.com/	George Miller
2	http://www.thedivergentseries.movie/#insurgent	Robert Schwentke

	tagline ... \
0	The park is open. ...
1	What a Lovely Day. ...
2	One Choice Can Destroy You ...

	overview runtime \
0	Twenty-two years after the events of Jurassic ... 124
1	An apocalyptic story set in the furthest reach... 120
2	Beatrice Prior must confront her inner demons ... 119

	genres \
0	Action Adventure Science Fiction Thriller
1	Action Adventure Science Fiction Thriller
2	Adventure Science Fiction Thriller

	production_companies	release_date
vote_count \		
0	Universal Studios Amblin Entertainment Legenda... 5562	6/9/15
1	Village Roadshow Pictures Kennedy Miller Produ... 6185	5/13/15
2	Summit Entertainment Mandeville Films Red Wago... 2480	3/18/15

	vote_average	release_year	budget_adj	revenue_adj
0	6.5	2015	1.379999e+08	1.392446e+09
1	7.1	2015	1.379999e+08	3.481613e+08
2	6.3	2015	1.012000e+08	2.716190e+08

[3 rows x 21 columns]

Observation

The revenue and budget have no currency specified, however, the TMDb from [themoviedb](#),
 ",All currency values on TMDb is assumed to be USD."

#dimension of dataframe

```
samples, features = movie_df.shape
print('The dataset has {} movie samples and {}
features'.format(samples, features))
```

The dataset has 10866 movie samples and 21 features

#summary statistics

```
movie_df.describe()
```

	id	popularity	budget	revenue
count	10866.000000	10866.000000	1.086600e+04	1.086600e+04
mean	66064.177434	0.646441	1.462570e+07	3.982332e+07
std	92130.136561	1.000185	3.091321e+07	1.170035e+08
min	5.000000	0.000065	0.000000e+00	0.000000e+00
25%	10596.250000	0.207583	0.000000e+00	0.000000e+00
50%	20669.000000	0.383856	0.000000e+00	0.000000e+00
75%	75610.000000	0.713817	1.500000e+07	2.400000e+07
max	417859.000000	32.985763	4.250000e+08	2.781506e+09

	vote_count	vote_average	release_year	budget_adj
count	10866.000000	10866.000000	10866.000000	1.086600e+04
mean	217.389748	5.974922	2001.322658	1.755104e+07
std	575.619058	0.935142	12.812941	3.430616e+07
min	10.000000	1.500000	1960.000000	0.000000e+00

0.000000e+00				
25%	17.000000	5.400000	1995.000000	0.000000e+00
0.000000e+00				
50%	38.000000	6.000000	2006.000000	0.000000e+00
0.000000e+00				
75%	145.750000	6.600000	2011.000000	2.085325e+07
3.369710e+07				
max	9767.000000	9.200000	2015.000000	4.250000e+08
2.827124e+09				

some deductions from the summary decription

- Most movies turn out to be profitable, althou some m
- Most movies have an average runtime of 102 mins
- The dataset contains no details of movies released after 2015 ##### Note:

This is a summary descriptive statistics, and is not a subtitue for the required indepth analysis and conclusion

Data Wrangling

1. Create a new dataframe(df) from a copy of the Movie_df

2. Drop duplicate movie samples

3. Add features relevant to the analysis

we create a new dataframe using the columns relevant to our analysis

```
['original_title', 'runtime', 'release_year', 'budget', 'revenue', 'budget_adj', 'revenue_adj']
```

Other features are important but not with respect to our analysis

1. create new dataframe from our old dataframe

```
cols_to_use =
np.array(['original_title', 'runtime', 'release_year', 'budget', 'revenue', 'budget_adj', 'revenue_adj'])
df = movie_df[cols_to_use].copy()
df.head()
```

	original_title	runtime	release_year	budget
revenue \ 0	Jurassic World	124	2015	150000000
1513528810				
1	Mad Max: Fury Road	120	2015	150000000
378436354				
2	Insurgent	119	2015	110000000
295238201				
3	Star Wars: The Force Awakens	136	2015	200000000
2068178225				

4	Furious 7	137	2015	190000000
1506249360				

	budget_adj	revenue_adj
0	1.379999e+08	1.392446e+09
1	1.379999e+08	3.481613e+08
2	1.012000e+08	2.716190e+08
3	1.839999e+08	1.902723e+09
4	1.747999e+08	1.385749e+09

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10866 entries, 0 to 10865
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   original_title  10866 non-null  object
1   runtime         10866 non-null  int64
2   release_year    10866 non-null  int64
3   budget          10866 non-null  int64
4   revenue         10866 non-null  int64
5   budget_adj      10866 non-null  float64
6   revenue_adj     10866 non-null  float64
dtypes: float64(2), int64(4), object(1)
memory usage: 594.4+ KB
```

clearly there are no empty cells in our new dataframe

2. Drop duplicate movie samples

```
print("The new tmdb dataset has {} duplicated movie
samples".format(df.duplicated().sum()))
```

The new tmdb dataset has 1 duplicated movie samples

We then drop duplicate rows(samples) keeping the first instance of their occurrence.

```
df.drop_duplicates(keep='first', inplace=True)
```

Check for features with null values and since their datatypes are not int/float, we ignore them.

```
print("the runtime feature has {} nan
values".format(df['runtime'].isna().sum()))
```

the runtime feature has 0 nan values

3. Add feature(s) relevant to our analysis.

```
df.insert(5, 'profit', df['revenue']-df['budget'])
df.insert(8, 'profit_adj', df['revenue_adj']-df['budget_adj'])
df.head()
```

	original_title	runtime	release_year	budget
revenue \				

0	Jurassic World	124	2015	150000000
1513528810				
1	Mad Max: Fury Road	120	2015	150000000
378436354				
2	Insurgent	119	2015	110000000
295238201				
3	Star Wars: The Force Awakens	136	2015	200000000
2068178225				
4	Furious 7	137	2015	190000000
1506249360				

	profit	budget_adj	revenue_adj	profit_adj
0	1363528810	1.379999e+08	1.392446e+09	1.254446e+09
1	228436354	1.379999e+08	3.481613e+08	2.101614e+08
2	185238201	1.012000e+08	2.716190e+08	1.704191e+08
3	1868178225	1.839999e+08	1.902723e+09	1.718723e+09
4	1316249360	1.747999e+08	1.385749e+09	1.210949e+09

Exploratory Data Analysis

function to compute data and make our plots

```
def feature_info_plot(col_name):
    #data_setup
    Title = f"yearly {col_name}"
    temp_data = df.groupby('release_year')[col_name].sum()

    #plot_face
    plt.figure(figsize=(15,5))
    plt.plot(temp_data)
    plt.title(Title)
    plt.xlabel('year')
    plt.ylabel(col_name)
    plt.show()
```

function for max and min values

```
def high_and_low(col_name):
    #for highest earned profit
    high= df[col_name].idxmax()
    high_details=pd.DataFrame(df.loc[high])

    #for lowest earned profit
    low= df[col_name].idxmin()
    low_details=pd.DataFrame(df.loc[low])

    #collectin data in one place
    info=pd.concat([high_details, low_details], axis=1)

    return info
```


function for line of best fit

```
def plot_best_fit(col_name):  
    #data setup  
    Title = f"yearly {col_name}"  
    temp_data = df.groupby('release_year')[col_name].sum()  
    pds = pd.Series(temp_data)  
    years = np.array(pds.index.values)  
    yearly_values = np.array([pds[i] for i in years])  
    x,y = years, yearly_values  
  
    #gradient and intercept  
    m,c = np.polyfit(x, y, 1)  
  
    #plot of best fit  
    plt.figure(figsize=(15,5))  
    plt.plot(x,y,'*')  
    plt.plot(x, m*x + c)  
    plt.title(Title)  
    plt.xlabel('year')  
    plt.ylabel(col_name)  
    plt.show()
```


Q1. How has the cost of movie production changed over the years?

Note: The Budget is the overhead cost of movie production.

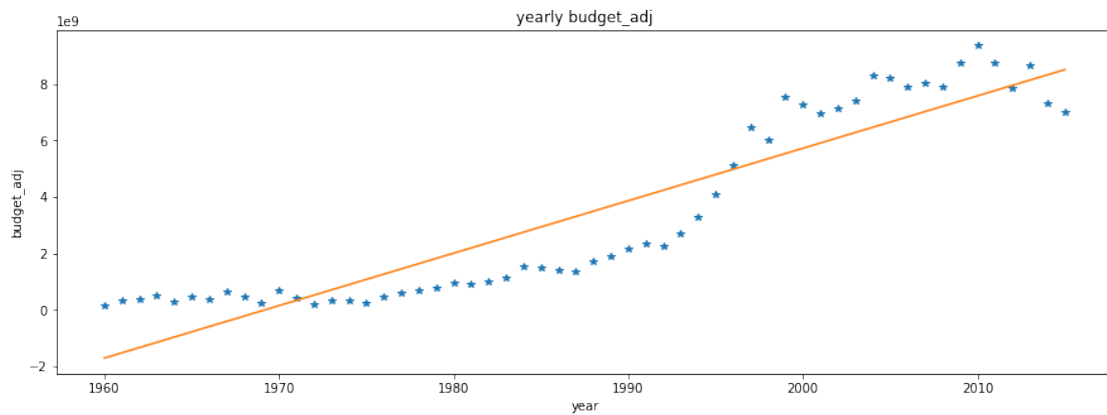
```
#plot for yearly budget  
feature_info_plot('budget')
```



```
#plot for yearly budget adjusted to the 2010 Dollar value  
feature_info_plot('budget_adj')
```



```
#budget plot of best fit
plot_best_fit('budget_adj')
```



```
pd.concat([high_and_low('budget'), high_and_low('budget_adj')],
axis=1)
```

	2244	30	2244
\			
original_title	The Warrior's Way	Mr. Holmes	The Warrior's Way
runtime	100	103	100
release_year	2010	2015	2010
budget	425000000	0	425000000
revenue	11087569	29355203	11087569
profit	-413912431	29355203	-413912431
budget_adj	425000000.0	0.0	425000000.0
revenue_adj	11087569.0	27006774.877019	11087569.0

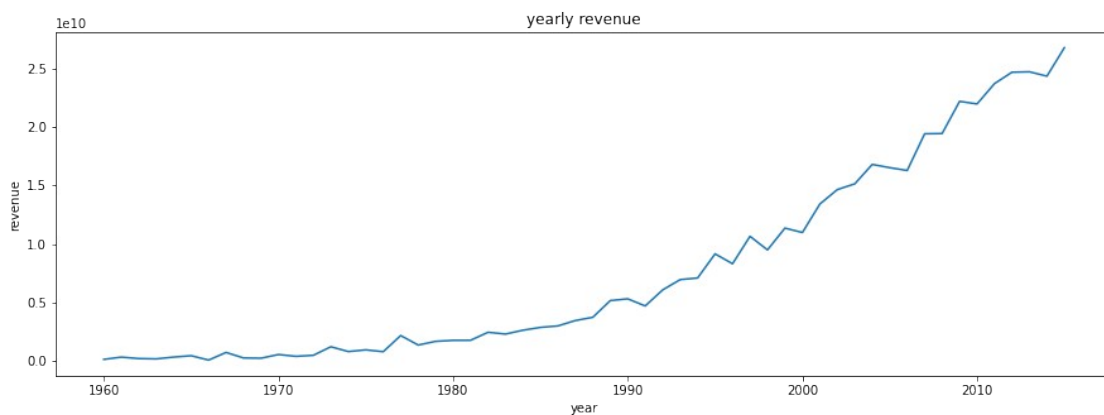
profit_adj -413912431.0 27006774.877019 -413912431.0

 30
original_title Mr. Holmes
runtime 103
release_year 2015
budget 0
revenue 29355203
profit 29355203
budget_adj 0.0
revenue_adj 27006774.877019
profit_adj 27006774.877019

Q2. How has revenue from movie sales and publicity changed over the the years


#plot for yearly revenue unadjusted

`feature_info_plot('revenue')`




#plot for yearly revenue adjusted

`feature_info_plot('revenue_adj')`



`plot_best_fit('revenue_adj')`



```
pd.concat([high_and_low('revenue'), high_and_low('revenue_adj')],
axis=1)
```


	1386	48	1386
\			
original_title	Avatar	Wild Card	Avatar
runtime	162	92	162
release_year	2009	2015	2009
budget	237000000	30000000	237000000
revenue	2781505847	0	2781505847
profit	2544505847	-30000000	2544505847
budget_adj	240886902.887613	27599987.856005	240886902.887613
revenue_adj	2827123750.41189	0.0	2827123750.41189
profit_adj	2586236847.524277	-27599987.856005	2586236847.524277

	48
original_title	Wild Card
runtime	92
release_year	2015
budget	30000000
revenue	0
profit	-30000000
budget_adj	27599987.856005
revenue_adj	0.0
profit_adj	-27599987.856005

Q3. How Profitable has movie production been over the years.

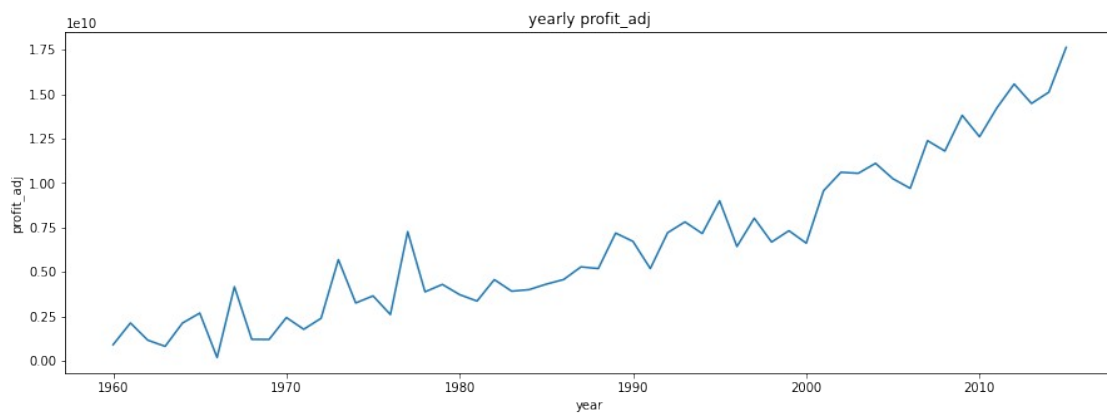
#plot for profit values unadjusted

```
feature_info_plot('profit')
```




#plot for adjusted profit values

```
feature_info_plot('profit_adj')
```



#profit plot of best fit

```
plot_best_fit('profit_adj')
```



```
pd.concat([high_and_low('profit'), high_and_low('profit_adj')],  
axis=1)
```

	1386	2244	
1329 \			
original_title	Avatar	The Warrior's Way	Star Wars
runtime	162	100	
121			
release_year	2009	2010	
1977			
budget	237000000	425000000	
11000000			
revenue	2781505847	11087569	
775398007			
profit	2544505847	-413912431	
764398007			
budget_adj	240886902.887613	425000000.0	
39575591.358274			
revenue_adj	2827123750.41189	11087569.0	
2789712242.27745			
profit_adj	2586236847.524277	-413912431.0	
2750136650.919176			

	2244
original_title	The Warrior's Way
runtime	100
release_year	2010
budget	425000000
revenue	11087569
profit	-413912431
budget_adj	425000000.0
revenue_adj	11087569.0
profit_adj	-413912431.0

Conclusions

1. Budget

- There was a sharp increase in cost of movie production from the early 1990s and peaked in the year 2009-2010 and then begins to decline.
- The rising cost of movie production between the 1990s and the year 2010 could be attributed to the use of state of the art technological equipments in movie production.
- The decline from the year 2010 could be attributed to increased adoption which has cut down the required manpower, thereby reducing the overall budget.

2. Revenue

- Movie revenue has seen an almost steady increase since the birth of movie production
- Inflation has had little effect on Movie Revenue.
- This shows growing interest in viewership.

3. Profit

- The overall profit shows some seasonality but has a summarily profitable trend
- from the high and low dataframes, we can see that a movie budget is not always an indicator high revenue or profit

Financial advice

- For an investor who wants to know if an investments in the movie industry is a worthy in the longterm, my answer is this since there is no significant difference in the adjusted and axctual values of the budget, revenue, and profit, the data at hand that the movie Industry is investment worthy.